

COMPETITIVE COVERAGE ANALYSIS

Florida Blue **vs.** Oscar Health Medical & Drug Coverage Policies


An AI-assisted, criterion-by-criterion comparison of 1,153 published coverage policies — where the two payers agree, where their requirements diverge, and which payer runs the tighter utilization-management criteria.

Prepared by **Mike Zehrer** · June 2026

Interactive site: payer-policy-cmp-06222027.web.app

Source & methodology: github.com/Elephruit/Medical-Policy

AT A GLANCE



KEY FINDINGS

01 Oscar runs materially tighter utilization management


Across the 126 overlapping drugs and services an LLM read on both sides, Oscar imposes the stricter coverage criteria on 92 (73%) of topics versus Florida Blue's 26 (21%); the rest are comparable. The asymmetry holds at the high end: 43 topics are *substantially* tighter on Oscar versus 13 on Florida Blue. Read as a business signal, Oscar's posture likely suppresses utilization and cost — at a higher risk of member and provider abrasion.




02 The two payers tighten in different ways

Counting the requirements each payer imposes that the other doesn't, the *shape* of each program is distinct. Oscar leans on **access gates** — step therapy, age limits, and specialist-prescriber requirements. Florida Blue leans on **operational controls** — dose/quantity limits and continuation-of-therapy rules — with far fewer step-therapy or specialist gates. Same goal, different levers: Oscar gatekeeps *who and when*; Florida Blue controls *how much and how long*.

Oscar adds most often



Florida Blue adds most often



03 Both programs share one backbone: carve-outs and continuation rules

The single most common 'extra' requirement on *both* payers is explicit **exclusions / non-covered uses** (0 on Oscar, 0 on Florida Blue), and **continuation-of-therapy rules** appear about equally on both (0 vs 0). Where they differ is everything layered on top of that shared base.

04 **AI subject-normalization recovered matches rule-based matching missed**

Matching policies across payers by document title alone is brittle — the same drug is titled differently by each insurer. Normalizing every one of the **1,153** policies to a canonical subject with an LLM surfaced **11** additional cross-payer matches that the deterministic title matcher could not link — same drug, different paperwork.

05 **“Coverage gaps” are mostly organizational, not real holes**

Florida Blue publishes **599** dedicated guidelines with no Oscar counterpart, versus **163** the other way. But much of that gap is *structure*, not coverage: Oscar consolidates whole drug classes into a single guideline — **8** class guidelines stand in for **34** separate Florida Blue per-drug policies. A genuine gap analysis has to read content, not count documents.

06 The sharpest single-drug divergences

The topics where the two payers' criteria differ most — the first places a reviewer should look. Each summary below is the model's one-line read of the key difference.

Ravulizumab (Ultomiris) IV or Subcutaneous

OSCAR MODERATE

Oscar is more restrictive overall: it covers NMOSD (which Florida Blue does not address), requires a minimum age of 18 for gMG and NMOSD, mandates failure of at least TWO standard therapies for gMG (vs. Florida Blue's ONE immunosuppressant), and imposes detailed TMA laboratory confirmation criteria for aHUS and broader PNH disease-burden thresholds (including reticulocyte count, renal dysfunction, pulmonary hypertension, dysphagia). Florida Blue is more granular with weight-based dosing limits for each indication and requires meningococcal vaccination documentation and a "no active infection" attestation that Oscar does not explicitly list.

Autonomic Nervous System Testing

OSCAR SUBSTANTIAL

Oscar's criteria are substantially more restrictive: rather than Florida Blue's broad three-part gateway (symptoms present, diagnosis unclear, result changes management), Oscar requires highly specific clinical prerequisites per test type—detailed syncope workup exclusions, contraindication screening, prior cardiac and medication evaluations, and an enumerated approved-diagnosis list for sudomotor testing. Florida Blue applies a single unified set of principles across all ANS testing, while Oscar creates separate, granular criteria for tilt-table and sudomotor testing.

Lyfgenia

OSCAR MODERATE

Oscar is moderately more restrictive than Florida Blue: it adds a performance-status gate (Karnofsky/Lansky $\geq 60\%$), a cerebral vasculopathy exclusion, a bone-marrow function floor, hepatitis B/C serologic requirements, a backup CD34+ rescue-cell collection mandate, an explicit CD34+ dose ceiling, a DMSO/dextran-40 hypersensitivity exclusion, and a specialist-prescriber requirement — none of which Florida Blue imposes. Florida Blue's sole unique additions are a tighter VOC threshold option (2+ events in 12 months) and a 9-month approval window vs. Oscar's 6-month window.

Pegcetacoplan (Empaveli) Infusion

FLORIDA BLUE SUBSTANTIAL

These two policies cover entirely different indications for pegcetacoplan: Florida Blue covers Empaveli (subcutaneous infusion) for PNH and C3G/IC-MPGN, while Oscar covers Syfovre (intravitreal injection) for geographic atrophy secondary to AMD — there is essentially zero clinical overlap in the criteria being compared, making a meaningful side-by-side alignment impossible.

Dupilumab (Dupixent) Injection

OSCAR SUBSTANTIAL

Oscar is substantially more restrictive than Florida Blue across nearly every indication: it mandates specific prior-therapy failure thresholds (e.g., ≥ 2 exacerbations, eosinophil counts, BSA criteria, OCS trials), explicit age floors, and defined approval durations, whereas Florida Blue primarily requires current/ongoing use of standard-of-care therapies rather than documented failures of them.

FINDINGS FROM READING THE POLICIES

Oscar gates drugs on prescriber specialty; Florida Blue usually doesn't

For specialty drugs, Oscar routinely requires the prescriber to be (or consult) a relevant specialist. Florida Blue's equivalent policies generally omit a specialist requirement, leaning on indication tables and FDA labeling.

Step therapy is codified differently

Oscar frequently requires a documented trial-and-failure of specific prior therapies; Florida Blue's structure centers on continuation-of-therapy and discloses Florida's step-therapy exception law.

Opposite consolidation strategies — and Oscar delegates some domains

Florida Blue consolidates drug classes into combined guidelines; Oscar keeps one guideline per drug but bundles some medical/biologic classes — and routes domains like advanced imaging to eviCore.

Bariatric surgery: Oscar maintains a dedicated adolescent policy

Oscar splits bariatric surgery into adult (CG008) and adolescent (CG009) guidelines with pediatric BMI-percentile criteria; Florida Blue addresses it within a single policy.

Catalog scope: 124 overlapping topics, hundreds unique — mostly structural

Florida Blue publishes far more individual guidelines (754 vs 387). Of these, 124 topics match across payers; 607 are Florida-Blue-only and 170 Oscar-only — but most “only” topics reflect granularity, class-bundling or vendor delegation, not proven non-coverage.

Oscar sets explicit age/weight minimums; Florida Blue defers to FDA labeling

Oscar's drug criteria commonly hard-code a minimum age (and sometimes weight). Florida Blue more often ties age to “within FDA labeling,” leaving it indication-dependent.

Florida Blue restricts site of care toward self-administration

Several Florida Blue drug policies declare that provider-administered (office/outpatient) administration is “not medically necessary” for self-injectable products — a site-of-care limit Oscar's matching policies don't emphasize.

SMA gene therapy (Zolgensma): different eligibility ceilings

Both require genetic confirmation, but Oscar caps eligibility at age <2 and disease onset before 6 months, while Florida Blue instead leans on SMN2 copy number, anti-AAV9 antibody titer and ventilation status.

Strong convergence on objective-threshold topics

Where coverage hinges on objective lab/imaging thresholds, the two payers largely agree — home oxygen and major-joint arthroplasty are near-identical.

HOW THIS WAS BUILT

Every policy was scraped from each payer's public clinical-guideline site, parsed from PDF to text, then **normalized by an LLM** into a canonical subject so the same drug or service matches across payers even when document titles differ. Overlapping topics were then **aligned and scored for restrictiveness by an LLM** reading both policies' coverage criteria. Restrictiveness is a decision-support signal, not ground truth — every claim links back to the source criteria on the interactive site. A two-tier model strategy (a fast model for the 1,153 per-policy

normalizations, a stronger model for the nuanced comparisons) keeps the whole enrichment a bounded, cached, reproducible cost.

```
scrape → extract (PDF→text) → LLM normalize (canonical subject) → match  
      → LLM compare + restrictiveness score → static site + this report
```

APPENDIX · METHODOLOGY & TECHNICAL FLOW

The full pipeline is open source at github.com/Elephruit/Medical-Policy. Each stage is a separate, re-runnable command; every LLM call is cached on disk by a content hash, so the analysis is incremental and reproducible.

1 Acquire

A per-payer `SourceAdapter` yields a catalog and fetches each policy PDF. Florida Blue is a stateful ASP.NET/Telerik site crawled by replaying postbacks (sequential); Oscar is a Next.js/Contentful site whose PDFs live on a CDN (stateless, parallel). Documents land in SQLite with a content hash for change detection.

2 Extract

PDF bytes → full text plus structured fields (policy number, authoritative subject, effective/revision dates, CPT/HCPCS codes). Source-agnostic.

3 Normalize every policy LLM · fast model

One call per policy distills it to a **canonical subject** (generic drug INN or standard service name), brand names, type, and key requirements — a payer-agnostic identity key. This is what lets the same drug match across insurers when their document titles differ.

4 Match into cross-payer topics

A union-find combines two signals: deterministic IDF-weighted cosine over title tokens (token-blocked, byte-stable), plus conservative LLM same-subject links from stage 3. A topic is flagged *AI-matched* only when it is cross-payer *solely* because of the LLM links — verified by a baseline diff against the lexical matcher.

5 Compare & score restrictiveness LLM · stronger model

For each overlapping topic, one call reads both payers' criteria and returns an aligned comparison — shared requirements (flagged same/differs), requirements unique to each payer, and a restrictiveness verdict (which payer is harder to get approved under, with rationale and a cost-vs-abrasion note).

6 Serve

Everything exports to a static JSON bundle behind a React/Vite site on Firebase Hosting — no database, no server — and to this report.

Two-tier model strategy

A fast model for the 1,153 high-volume normalizations; a stronger model for the ~127 nuanced comparisons. Bounded, cached, one-time cost.

Structured output

Forced tool-use pins each call to a single schema, so every result is valid JSON without bespoke parsing.

Determinism & idempotence

The lexical matcher is byte-stable run-to-run; LLM results cache by (prompt version, model, inputs) — re-runs touch only changed inputs.

Decision support, not ground truth

Restrictiveness is an LLM judgment over PDF-extracted text; every claim links back to the source criteria on the interactive site.

